

*Induction: A process of reasoning (arguing) which infers a general conclusion based on individual cases*

# Supervised (Inductive) Learning

The University of Texas at Dallas

# Supervised Learning

- **Given:** Training examples  $\langle \mathbf{x}, f(\mathbf{x}) \rangle$  for some unknown function  $f$ .
- **Find:** A good approximation to  $f$ .

## Example Applications

- **Credit risk assessment**  
 $\mathbf{x}$ : Properties of customer and proposed purchase.  
 $f(\mathbf{x})$ : Approve purchase or not.
- **Disease diagnosis**  
 $\mathbf{x}$ : Properties of patient (symptoms, lab tests)  
 $f(\mathbf{x})$ : Disease (or maybe, recommended therapy)
- **Face recognition**  
 $\mathbf{x}$ : Bitmap picture of person's face  
 $f(\mathbf{x})$ : Name of the person.
- **Automatic Steering**  
 $\mathbf{x}$ : Bitmap picture of road surface in front of car.  
 $f(\mathbf{x})$ : Degrees to turn the steering wheel.

## Appropriate Applications for Supervised Learning

- **Situations where there is no human expert**

$\mathbf{x}$ : Bond graph for a new molecule.

$f(\mathbf{x})$ : Predicted binding strength to AIDS protease molecule.

- **Situations where humans can perform the task but can't describe how they do it.**

$\mathbf{x}$ : Bitmap picture of hand-written character

$f(\mathbf{x})$ : Ascii code of the character

- **Situations where the desired function is changing frequently**

$\mathbf{x}$ : Description of stock prices and trades for last 10 days.

$f(\mathbf{x})$ : Recommended stock transactions

- **Situations where each user needs a customized function  $f$**

$\mathbf{x}$ : Incoming email message.

$f(\mathbf{x})$ : Importance score for presenting to user (or deleting without presenting).

# A learning problem!

X	0	X
0	X	0
0	X	X

X	0	X
X	X	0
X	0	0

X	X	X
0	X	X
0	0	0

$f(x)=1$

---

0	X	0
X	0	X
0	X	X

0	0	X
X	X	0
0	X	X

0	X	X
X	0	0
0	X	X

$f(x)=0$

---

0	X	X
0	X	0
X	X	0


$f(x)=?$

# If you prefer the training data in this form!

X1	X2	X3	X4	X5	X6	X7	X8	X9	f(x)
X	0	X	0	X	0	0	X	X	1
X	0	X	X	X	0	X	0	0	1
X	X	X	0	X	X	0	0	0	1
0	X	0	X	0	X	0	X	X	0
0	0	X	X	X	0	0	X	X	0
0	X	X	X	0	0	0	X	X	0
0	X	X	0	X	0	X	X	0	?

- $x$ : a 9-dimensional vector
- $f(x)$ : a function or a program that takes the vector as input and outputs either a 0 or a 1
- **Task**: given the training examples, find a good approximation to  $f$  so that in future if you see an unseen vector “ $x$ ” you will be able to figure out the value of  $f(x)$

## A Learning Problem



**A simpler  
example for  
analysis!**

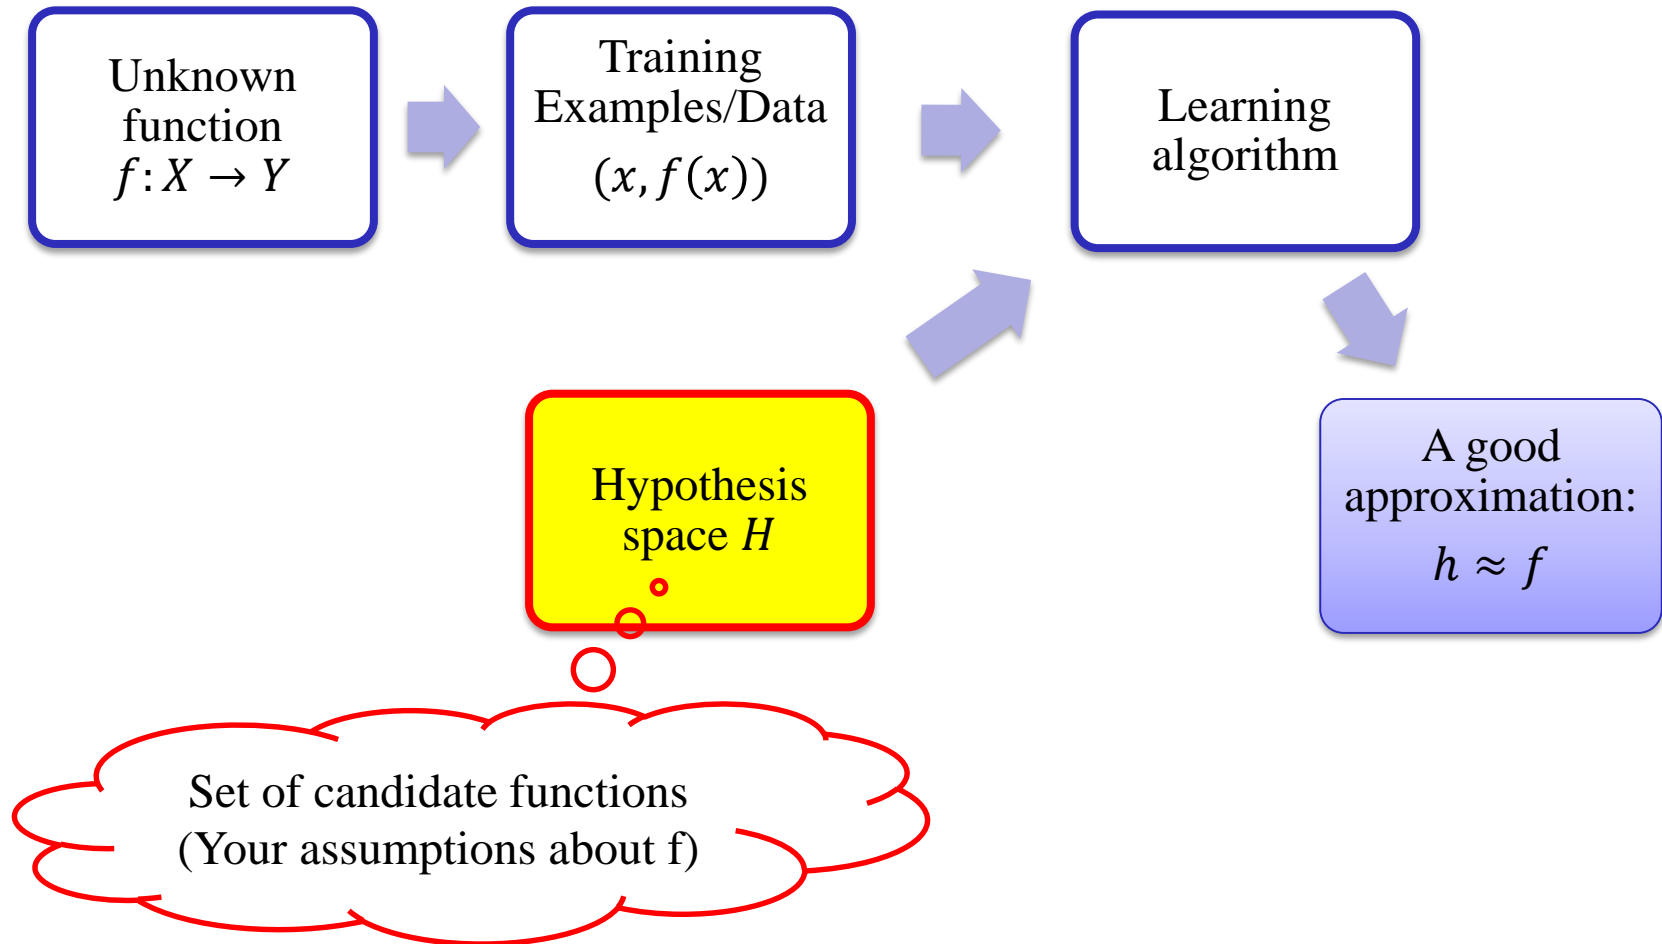
Example	$x_1$	$x_2$	$x_3$	$x_4$	$y$
1	0	0	1	0	0
2	0	1	0	0	0
3	0	0	1	1	1
4	1	0	0	1	1
5	0	1	1	0	0
6	1	1	0	0	0
7	0	1	0	1	0

Classification problem

Given data or examples, find the function  $f$ ?

# How to find a good approximation to $f$ ?

- A possible/plausible technique



## Hypothesis Spaces

- **Complete Ignorance.** There are  $2^{16} = 65536$  possible boolean functions over four input features. We can't figure out which one is correct until we've seen every possible input-output pair. After 7 examples, we still have  $2^9$  possibilities.

You are assuming that the unknown function  $f$  could be any one of the  $2^{16}$  functions!

$x_1$	$x_2$	$x_3$	$x_4$	$y$
0	0	0	0	?
0	0	0	1	?
0	0	1	0	0
0	0	1	1	1
0	1	0	0	0
0	1	0	1	0
0	1	1	0	0
0	1	1	1	?
1	0	0	0	?
1	0	0	1	1
1	0	1	0	?
1	0	1	1	?
1	1	0	0	0
1	1	0	1	?
1	1	1	0	?
1	1	1	1	?

It turns out that out of the  $2^{16}$  possible functions,  $2^9$  classify all points in the training data correctly!



## Hypothesis Spaces (2)

- **Simple Rules.** There are only 16 simple conjunctive rules.

You are assuming that the unknown function  $f$  could be any one of the 16 conjunctive rules!

Unfortunately, none of them work

Rule	Counterexample
$\Rightarrow y$	1
$x_1 \Rightarrow y$	3
$x_2 \Rightarrow y$	2
$x_3 \Rightarrow y$	1
$x_4 \Rightarrow y$	7
$x_1 \wedge x_2 \Rightarrow y$	3
$x_1 \wedge x_3 \Rightarrow y$	3
$x_1 \wedge x_4 \Rightarrow y$	3
$x_2 \wedge x_3 \Rightarrow y$	3
$x_2 \wedge x_4 \Rightarrow y$	3
$x_3 \wedge x_4 \Rightarrow y$	4
$x_1 \wedge x_2 \wedge x_3 \Rightarrow y$	3
$x_1 \wedge x_2 \wedge x_4 \Rightarrow y$	3
$x_1 \wedge x_3 \wedge x_4 \Rightarrow y$	3
$x_2 \wedge x_3 \wedge x_4 \Rightarrow y$	3
$x_1 \wedge x_2 \wedge x_3 \wedge x_4 \Rightarrow y$	3

Example	$x_1$	$x_2$	$x_3$	$x_4$	$y$
1	0	0	1	0	0
2	0	1	0	0	0
3	0	0	1	1	1
4	1	0	0	1	1
5	0	1	1	0	0
6	1	1	0	0	0
7	0	1	0	1	0

No simple rule explains the data. The same is true for simple clauses.

## Hypothesis Space (3)

- ***m*-of-*n* rules.** There are 32 possible rules (includes simple conjunctions and clauses).

At least  $m$  of the  $n$   
variables must be true

You are assuming  
that the unknown  
function  $f$  could  
be any one of the  
32 *m*-of-*n* rules!

Only one of them, the  
one marked by “\*\*\*”  
works!

variables	Counterexample			
	1-of	2-of	3-of	4-of
$\{x_1\}$	3	—	—	—
$\{x_2\}$	2	—	—	—
$\{x_3\}$	1	—	—	—
$\{x_4\}$	7	—	—	—
$\{x_1, x_2\}$	3	3	—	—
$\{x_1, x_3\}$	4	3	—	—
$\{x_1, x_4\}$	6	3	—	—
$\{x_2, x_3\}$	2	3	—	—
$\{x_2, x_4\}$	2	3	—	—
$\{x_3, x_4\}$	4	4	—	—
$\{x_1, x_2, x_3\}$	1	3	3	—
$\{x_1, x_2, x_4\}$	2	3	3	—
$\{x_1, x_3, x_4\}$	1	***	3	—
$\{x_2, x_3, x_4\}$	1	5	3	—
$\{x_1, x_2, x_3, x_4\}$	1	5	3	3

Example	$x_1$	$x_2$	$x_3$	$x_4$	$y$
1	0	0	1	0	0
2	0	1	0	0	0
3	0	0	1	1	1
4	1	0	0	1	1
5	0	1	1	0	0
6	1	1	0	0	0
7	0	1	0	1	0

## Two Views of Learning

- **Learning is the removal of our remaining uncertainty.** Suppose we *knew* that the unknown function was an  $m$ -of- $n$  boolean function, then we could use the training examples to infer which function it is.
- **Learning requires guessing a good, small hypothesis class.** We can start with a very small class and enlarge it until it contains an hypothesis that fits the data.

**We could be wrong!**

- **Our prior knowledge might be wrong**
- **Our guess of the hypothesis class could be wrong**

The smaller the hypothesis class, the more likely we are wrong.

Example:  $x_4 \wedge \text{Oneof}\{x_1, x_3\} \Rightarrow y$  is also consistent with the training data.

Example:  $x_4 \wedge \neg x_2 \Rightarrow y$  is also consistent with the training data.

If either of these is the unknown function, then we will make errors when we are given new  $x$  values.

## Two Strategies for Machine Learning

- **Develop Languages for Expressing Prior Knowledge:** Rule grammars and stochastic models.
- **Develop Flexible Hypothesis Spaces:** Nested collections of hypotheses.  
Decision trees, rules, neural networks, cases.

In either case:

- **Develop Algorithms for Finding an Hypothesis that Fits the Data**

## Terminology

- **Training example.** An example of the form  $\langle \mathbf{x}, f(\mathbf{x}) \rangle$ .
- **Target function (target concept).** The true function  $f$ .
- **Hypothesis.** A proposed function  $h$  believed to be similar to  $f$ .
- **Concept.** A boolean function. Examples for which  $f(\mathbf{x}) = 1$  are called **positive examples** or **positive instances** of the concept. Examples for which  $f(\mathbf{x}) = 0$  are called **negative examples** or **negative instances**.
- **Classifier.** A discrete-valued function. The possible values  $f(\mathbf{x}) \in \{1, \dots, K\}$  are called the **classes** or **class labels**.
- **Hypothesis Space.** The space of all hypotheses that can, in principle, be output by a learning algorithm.
- **Version Space.** The space of all hypotheses in the hypothesis space that have not yet been ruled out by a training example.

# Key Issues in Machine Learning

- **What are good hypothesis spaces?**

Which spaces have been useful in practical applications and why?

- **What algorithms can work with these spaces?**

Are there general design principles for machine learning algorithms?

- **How can we optimize accuracy on future data points?**

This is sometimes called the “problem of overfitting”.

- **How can we have confidence in the results?**

How much training data is required to find accurate hypotheses? (the *statistical question*)

- **Are some learning problems computationally intractable?**

(the *computational question*)

- **How can we formulate application problems as machine learning problems?** (the *engineering question*)

# Steps in Supervised Learning

1. Determine the representation for “ $x, f(x)$ ” and determine what “ $x$ ” to use

## **Feature Engineering**

2. Gather a training set (not all data is kosher)

## **Data Cleaning**

3. Select a suitable evaluation method
4. Find a suitable learning algorithm among a plethora of available choices
  - Issues discussed on the previous slide

# Feature Engineering is the Key


- Most effort in ML projects is constructing features
- Black art: Intuition, creativity required
  - Understand properties of the task at hand
  - How the features interact with or limit the algorithm you are using.
- ML is an iterative process
  - Try different types of features, experiment with each and then decide which feature set/algorithm combination to use



# A sample machine learning Algorithm

- 2-way classification problem
  - +ve and –ve classes
- Representation: Lines ( $Ax+By=C$ )
  - Specifically
    - if  $Ax+By+C > 0$  then classify “+ve”
    - Else classify as “-ve”
- Evaluation: Number of mis-classified examples
- Optimization: An algorithm that searches for the three parameters: A, B and C.

# Toy Example



**Blue circles:**  
Good credit (low risk)

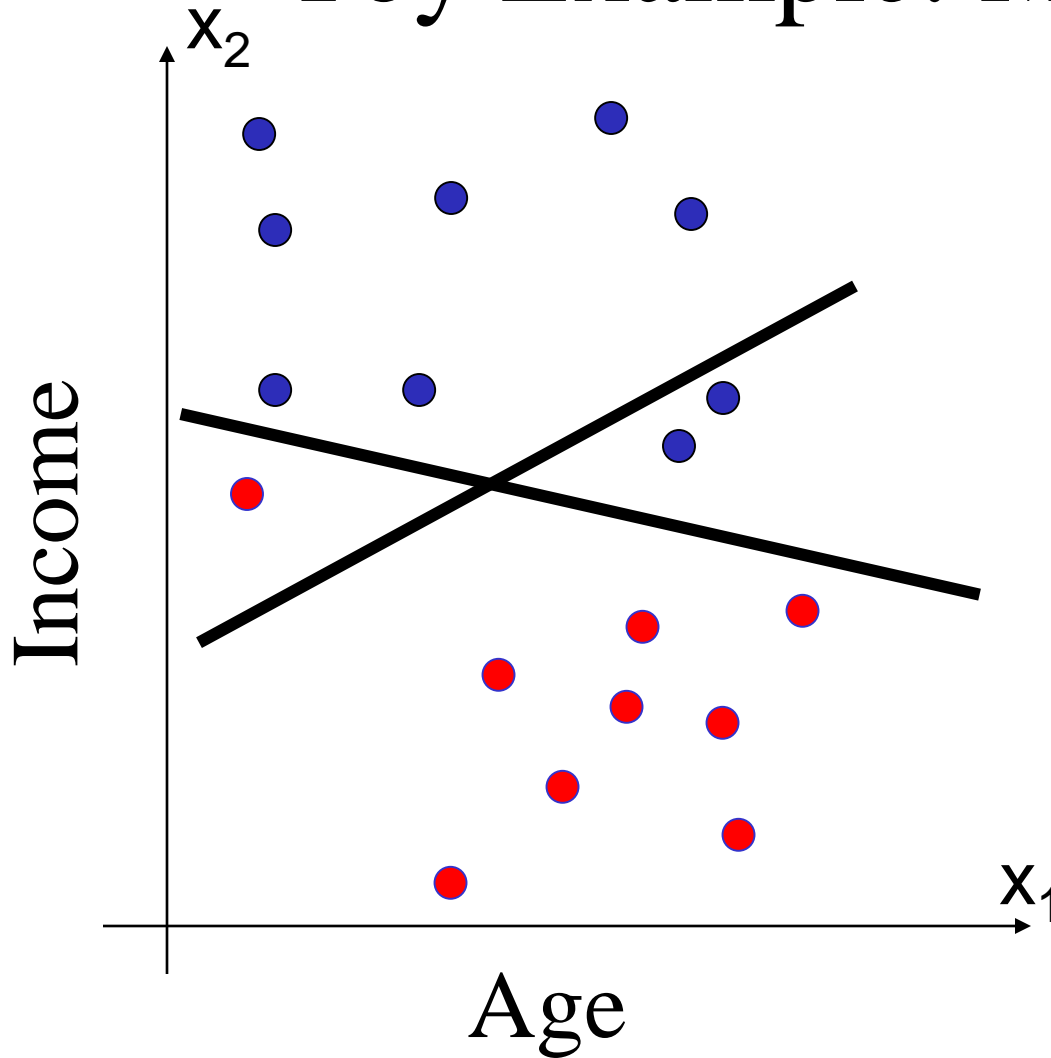
**Red circles:** Bad credit (high risk)

**Problem:** Fit a line that separates the two such that the error is minimized.

# How do machine learners solve this problem?

- Try different lines until you find one that separates the data into two
- A more plausible alternative
  - Begin with a random line
  - Repeat until no errors
  - For each point
    - If the current line says +ve and point is -ve then decrease A, B and C
    - If the current line says -ve and the point is +ve then increase A, B, and C

# Toy Example: More data



**Blue circles:** Good credit (low risk)

**Red circles:** Bad credit (high risk)

**Problem:** Fit a line that separates the two such that the error is minimized.

# Learning = Representation + Evaluation + Optimization

- Combinations of just three elements

Representation	Evaluation	Optimization
Instances	Accuracy	Greedy search
Hyperplanes	Precision/Recall	Branch & bound
Decision trees	Squared error	Gradient descent
Sets of rules	Likelihood	Quasi-Newton
Neural networks	Posterior prob.	Linear progr.
Graphical models	Margin	Quadratic progr.
Etc.	Etc.	Etc.