

A CHECKERS LEARNING PROBLEM

“Machine Learning” By Tom Mitchell

PROBLEM

- Task T: playing checkers
- Performance measure ***P***: percent of games won in the world tournament
- Training experience E: games played against itself

APPROACH

1. The exact type of knowledge to be learned
2. A representation for this target knowledge
3. A learning mechanism

The type of training experience available can have a significant impact on success or failure of the learner

TARGET FUNCTION

- to reduce the problem of improving performance P at task T to the problem of learning some particular ***target function***

Legal Moves	Won or Lost
-------------	-------------

← Indirect Training Experience

Given Broad	Moves required to win
-------------	-----------------------

← Direct Training Experience

TARGET FUNCTION

- Evaluation function that assigns a numerical score to any given board state.
 - $V : B \rightarrow \mathbb{R}$ to denote that V maps any legal board state from the set B to some real value (we use \mathbb{R} to denote the set of real numbers).
1. if b is a final board state that is won, then $V(b) = 100$
 2. if b is a final board state that is lost, then $V(b) = -100$
 3. if b is a final board state that is drawn, then $V(b) = 0$
 4. if b is not a final state in the game, then $V(b) = V(b')$, where b' is the best final board state that can be achieved starting from b and playing optimally until the end of the game (assuming the opponent plays optimally, as well).

TARGET FUNCTION

- ***operational description of the ideal target function*** V is required.
- Learning algorithms is expected to acquire only some ***approximation*** to the target function, and for this reason the process of learning the target function is often called ***function approximation***

On one hand, we wish to pick a very expressive representation to allow representing as close an approximation as possible to the ideal target function V . On the other hand, the more expressive the representation, the more training data the program will require in order to choose among the alternative hypotheses it can represent.

Problem Representation

A simple representation: for any given board state, the function \hat{V} will be calculated as a linear combination of the following board features.

- **$x1$** : the number of black pieces on the board
- **$x2$** : the number of red pieces on the board
- **$x3$** : the number of black kings on the board
- **$x4$** : the number of red kings on the board
- **$x5$** : the number of black pieces threatened by red (i.e., which can be captured on red's next turn)
- **$X6$** : the number of red pieces threatened by black

TARGET FUNCTION

Thus, our learning program will represent $\hat{V}(b)$ as a linear function of the form

$$\hat{V}(b) = w_0 + w_1x_1 + w_2x_2 + w_3x_3 + w_4x_4 + w_5x_5 + w_6x_6$$

where w_0 through w_6 are numerical coefficients, or weights, to be chosen by the learning algorithm. Learned values for the weights w_1 through w_6 will determine the relative importance of the various board features in determining the value of the board, whereas the weight w_0 will provide an additive constant to the board value.

ESTIMATING TRAINING VALUES

- In order to learn the target function \hat{V} we require a set of training examples, each describing a specific board state b and the training value $V_{train}(b)$ for b . In other words, each training example is an ordered pair of the form $\langle b, V_{train}(b) \rangle$.
- **Rule for estimating training values.**

$$V_{train}(b) \leftarrow \hat{V}(\text{Successor}(b))$$

ADJUSTING THE WEIGHTS

- One common approach is to define the best hypothesis, or set of weights, as that which minimizes the square error E between the training values and the values predicted by the hypothesis \hat{V} .
- $$E = \sum_{\langle b, V_{train}(b) \rangle \in \text{training sample}} (V_{train}(b) - \hat{V}(b))^2$$

Thus, we seek the weights, or equivalently the \hat{V} , that minimize E for the observed training examples.

LMS Training

Least mean squares or **LMS** training rule is one of several algorithms to incrementally refine the weights.

LMS weight update rule.

- **For each training example $\langle b, V_{train}(b) \rangle$**
 - **Use the current weights to calculate $\hat{V}(b)$**
 - **For each weight w_i , update it as**

$$w_i \leftarrow w_i + \eta (V_{train}(b) - \hat{V}(b)) x_i$$

The Final Design

